[TDDE09 2026] Lecture 2.3

Points: 3/5

✓ **Correct** 1/1 Points

1. Which of the following NLP tasks is most likely to be handled comparatively well using static word vectors rather than contextual embeddings?

- ○ coreference resolution
- ◉ topic classification
- ○ word sense disambiguation

✓ **Correct** 1/1 Points

2. Consider the sentence "Dogs may bark at strangers" and assume that words are indexed from 1 to 5. Which attention weight do you expect to be highest?

- ◉ α_{31}
- ○ α_{33}
- ○ α_{35}

✗ **Incorrect** 0/1 Points

3. Consider following values for the example of "Contextual word embeddings via attention".

h1 = [0.5539, 0.7239], h2 = [0.4111, 0.3878], h3 = [0.2376, 0.1264]

Assuming that the attention score is computed using the unscaled dot product, what is the refined representation for h2?

- (selected) [0.5084, 0.3194, 0.1467]
- [0.4198, 0.4488]
- [0.3962, 0.3279, 0.2759]

✕ **Incorrect** 0/1 Points

4. Which of the following statements about the more general characterisation of attention in terms of queries, keys and values is true?

- (selected) The query has the same length as each value.
- The output has the same length as each value.
- Each key has the same length as each value.

✓ **Correct** 1/1 Points

5. Consider an instance of multi-head attention with 8 heads where the queries and keys have size 256. What would be the typical key length in each block's attention mechanism?

- 256
- (selected) 32
- 8

Microsoft 365

This content is created by the owner of the form. The data you submit will be sent to the form owner. Microsoft is not responsible for the privacy or security practices of its customers, including those of this form owner. Never give

out your password.

Microsoft Forms | AI-Powered surveys, quizzes and polls Create my own form

Privacy and cookies | Terms of use