[TDDE09 2026] Lecture 2.4

Points: 4/5

✓ **Correct** 1/1 Points

1. Which of the following is the main advantage of the Transformer architecture over recurrent neural networks?

- ○ supports significantly more compact models
- ○ significantly reduced need for training data
- ◉ direct access to all elements in the input sequence

✗ **Incorrect** 0/1 Points

2. Consider the example translation used to illustrate the Transformer architecture. Which of the following statements is false?

- ○ The final encoder representation of *coffee* depends on the token embedding of *drink*.
- ◉ The final decoder representation of *Kaffee* depends on the final encoder representation of *coffee*.
- ○ The final encoder representation of *drink* depends on the token embedding of *Kaffee*.

✓ **Correct** 1/1 Points

3. The Transformer architecture uses three different variants of multi-head attention. Which one is used in the encoder?

- ○ cross-attention

- (●) self-attention
- ( ) masked self-attention

✓ **Correct** 1/1 Points

4. What is the purpose of layer normalisation?

- ( ) down-scaling the layer's output values
- ( ) squeezing the layer's input values into the interval [0, 1]
- (●) centering and scaling the layer's input values

✓ **Correct** 1/1 Points

5. True or false: Permuting the input tokens to a Transformer encoder does not change the final token representations.

- ( ) True
- (●) False
- ( ) Depends on the input tokens

Microsoft 365

This content is created by the owner of the form. The data you submit will be sent to the form owner. Microsoft is not responsible for the privacy or security practices of its customers, including those of this form owner. Never give out your password.

Microsoft Forms | AI-Powered surveys, quizzes and polls Create my own form

Privacy and cookies | Terms of use