[TDDE09 2026] Lecture 2.5

Points: 4/5

✓ **Correct** 1/1 Points

1. What does the term generative pre-training refer to?

- (●) pre-training on a language modelling task
- ( ) pre-training on automatically generated text
- ( ) pre-training with a generative probabilistic model

✗ **Incorrect** 0/1 Points

2. Looking at the original GPT model architecture (Radford et al., 2018), what is the approximate number of trainable parameters in the FNN?

- (●) 589824
- ( ) 9216
- ( ) 4718592

✓ **Correct** 1/1 Points

3. Suppose you want to fine-tune a GPT model on the Stanford Natural Language Inference dataset. What is the minimal number of parameters you need to update?

- ( ) the number of parameters in the pre-trained GPT model

- [x] the number of parameters in the final Linear layer
- [ ] the sum of these two

✓ **Correct** 1/1 Points

4. What do we mean when we say that GPT-3 exhibits zero-shot behaviour?

- [ ] It can solve tasks without receiving any input.
- [ ] It can solve tasks without any training.
- [x] It can solve tasks without any task-specific fine-tuning.

✓ **Correct** 1/1 Points

5. What is a resonable explanation for the observation that GPT-3 can translate from English to French?

- [x] The data sets used for pre-training contain example translations.
- [ ] The number of model parameters approaches the number of neurons in the human brain.
- [ ] The model has been trained based on feedback from professional translators.

Microsoft 365

This content is created by the owner of the form. The data you submit will be sent to the form owner. Microsoft is not responsible for the privacy or security practices of its customers, including those of this form owner. Never give out your password.

Microsoft Forms | AI-Powered surveys, quizzes and polls Create my own form

Privacy and cookies | Terms of use